Contents

SDN^2: Marketing System for Data-center Network Resource Allocation based on Software Defined Network

Rui Zhou　　　　　Shu Zhang

Final Project - CSCI 2950-u - Spring 2013

Abstract

In this report we present SDN^2, an auction-based Marketing System for Datacenter Network Resource Allocation based on Software Defined Networks. Motivated by similar mechanism of auctions for resource allocation in other computer systems, we aim to let users bid their network resources. Currently our focus mainly lays in bandwidth and latency for desired flows. Different from the only bidding system for Network Resource existing in Google's Data-center[9], our system is based on Software Defined Network and takes the advantage of the central controller and global network information. Another aspect of our focus is on the allocation algorithm for the auctioneer. We proposed and investigated some algorithms as well as evaluated the performances of some of them.

1 INTRODUCTION

Resource allocation is an important topic in system related area, especially in data centers where resources distribution heavily influences the performance of the application running in the data-centers. Since the network resource is limited and application expense could be infinite, the idea of launching auctions on data-centers for applications to bid resources comes to existence. With the core of auction, some marketing mechanisms have been built for public or private data-centers, such as AWS[1] or internal data-centers for large companies such as Google and Microsoft. While various types of marketing systems emerges, they primarily focus on the allocation of computation power and storage room, few of them cares enough about network resources, let along combining with the existence of SDN network.

The emergence of SDN network provides network manipulators with more powerful ways to manage their networks. For example, the existence of NIB(Network Information Base) in Onix[7] could provide the upper layer with the global topology of the entire network just in a single manager machine. SDN also provide the ability to configure the switches in the networks on the fly, thus a more dynamic way to bid and evolve the network could be built.

In this paper we combine the idea of auction in data-centers networking with the SDN technology so as to reach the goal of bidding network resources. Since we have the central information of the network, bidding and allocations of requested resources could be performed in a single node in the network(aka. the controller) quickly to reach the goal of real-time allocation. The general work flow is shown in figure 1.

Figure 1: Auction Working Flow

The general flow for each round of bidding starts from the bidders. Bidders first recognize their flow requirements in the near future, for example, they might need to reserve the resource for a flow from host A to host B in the network with a certain amount of data to transfer, with the hope of minimum bandwidth and latency (QoS) guarantees in a certain time period. After confirming these requirements, bidders send bid-

ding requests to the auctioneer. The auctioneer collects requests for this round and wait for the bidding clock to go off, saying it is time to calculate the allocation. With specific allocation algorithm, the auctioneer could get a Boolean result for each bidder's request and send the feedback to the bidders, and then a new round starts.

Our goals for this project are to build a workable system running the auction and doing allocation in real time, to provision a protocol (a standard interface) for all parties wishing to use the system to follow so as to communicate using the protocol, and to investigate algorithms to allocation resources basing on requests, for different goals.

In section 2 we investigate some existing backgrounds that serves as inspiration or basement for this project. In sections 4 we conclude the contribution of our work. In section 5 the design and algorithms of the system are discussed and in the implementation is introduced in section 6. In section 7 we do some experiments and evaluation on the out-coming products. Since our system is still in prototype, in section 7 we have identified some extensible work which would be done in the future. And in section 8 we discuss some related work.

2 KEYWORDS & CATEGORIES

Data-center, Network Resource Allocation, Marketing and Auction, Software Defined Network.

3 BACKGROUND

In this section, we investigate existing data centers (mainly private data-centers and public data-centers) and their marketing and pricing models. Also, we introduce the existing techniques of SDN which is the base of our system. These investigations essentially inspired our idea of this project.

3.1 Datacenter Markets

There are generally two types of data centers, one is for-profit, public cloud data-centers, such as Amazon Web Services[1], the other is private cloud data-centers, like those operated by Google or Facebook. Although auction mechanisms could be applied to both, different scenarios results in diverse answers because of their business models. Although most of the models applies in computation and storage marketing,

they serve the same principle for our Network Allocation System.

In a public cloud service, bidders are customers who use the cloud service. Take Amazon EC2 for example[3], Amazon revises prices for instances for all kinds of types. Three basic types are On-Demand Instances, Reserved Instances, which allows user to reserve a period of time for future uses and Spot Instances[4] whose values fluctuate and allow customers to bid on them. Along this path, Amazon created other services like data transfer, elastic IP addresses, elastic load balancing, etc and makes pricing policies for them. From Amazon's view, its pricing policies are so carefully revised so that they could maximize its profits. If a customer wants to bid for Amazon's resources, her knowledge is pretty limited, and Amazon hides the data center details from customers and only expose pricing history as hints for users to bid, and the budget of buyers are real currency.

In a private data-centers running inside IT companies, the customers are internal teams of companies who use the data-centers resources for the sake of their team and the companies. The operator's goal might be minimizing the cost while meeting business objectives, or may be efficient use of network resources. The users (teams), according to their priorities, are allocated different amounts of budget, which usually is virtual currency. Teams with more budget might use the data-centers resources in a higher share. The private data-centers might perform as a autonomy market, whose prices of goods are totally determined by the market or self-adapted by the system. Bidders, with higher possibility, will be provided more detailed view of the data-centers conditions, such as dynamic traffic between servers or racks and static network topology, so that they could formulate their bidding policies more sensible.

3.2 Existing Pricing Mechanisms

We found that the value of goods which are to be traded in the market is the core issue when letting a market run. People may argue that in an auction, people could win the bidding by giving the highest amount of money per unit of resource. But sometimes it might be far more complex. Prices, in some cases, are the reflection of the condition of the market, so it is better to set a baseline price for resources so that people who bid lower than that baseline will never

win the auction, even if they have the highest price. But in some cases, prices do not need to be considered and the simple first-price or second-price auctions policy could perform well. Also, complex cases like combinatorial pricing[10] is also a must-to-think issue. Also, The pricing methods for auctions in private and public data-centers are pretty different. This is because the goals of markets in either data-centers are different.

3.2.1 Public Cloud Pricing

Amazon[1] has set models for pricing in public cloud. AWS supports buy-at-once pricing for on-demand instances as well as auctions for resources for spot instances. For on-demand services, prices for instances are relatively higher and the good thing is that they are guaranteed to be provided to the buers. For Spot instances[4], Amazon has set the base prices for instances. The prices are much lower ($0.007 per hour versus $0.060 per hour for on-demand instances). The prices are given by Amazon according to their costs and virtual equipments with different capacities will have different prices. Resources are provided in a combinatorial way, since it is meaningless to bid a single CPU without any memory or storage. But in some scenarios resources are not sold in a combinational way, such as Elastic Load Balancing which only charges by the amount of data to be processed by the load balancer.

3.2.2 Private Cloud Pricing

Unlike AWS, prices per unit of resources in the private cloud data-centers probably will change in a dynamic way. Systems like D'Agents[12] and Google's Planet-wide Clusters[9] has adopted dynamic pricing and price of resources are updated periodically. Both the aforementioned systems uses dynamic prices to reflect the changing demand-supply relation. For dynamic pricing, the prices are the minimum payment. But in order to maximize the utility of total resources, the sealed-bid second-price auctions [13] could be used so as to guarantee that as long as there are more than one bidders in a certain auction round, there will always be a winner. The dynamic minimum prices in this case, serve as an indication for bidders to price their bid requests. The prices in [9] could change in an self-adaptive way. In short, 'hot' resources will be more expensive and a function $(g(x,p))$ calculates the price increment after each auction round, basing on

the exceeding of demands over supplies. Interestingly, AuYoung, et al [11] discovered that auctions might take some time as bidder's patience falls down when waiting for gaining the resource. So there should be another 'buy-it-now' pricing mechanism provided for users who don't want to wait. The buy-it-now pricing short-circuits the price discovery process of the auction, and the price is determined by a historical function which takes the historical auction/buy-it-now prices for parameters.

3.3 Software Defined Network

While the data-centers technology evolves rapidly, **Software Defined Network (SDN)** has also been proposed as a stunning idea. By separating the control panel and the data panel, SDN abstracts a supported network to be a networking system, which offers a central view of the networks as a graph and provides a central controller based on flows. Openflow is a dominating protocol used in SDN based Network systems which defines what is analog to the lower level API in an operating systems. And high level Networking systems such as NOX[22] and Floodlight[2], and Nettle[24] have been developed based on Openflow. SDN and Network systems has been actively developed in many universities. Companies such as Big Switch and JUNIPER have long started to build SDN supported switches. It is believed that the future of network belongs to SDN. In this paper, we also take SDN as the approach towards the effective allocation of Network resources in data-centers. Recently Andrew et al' proposed **Participatory Networking(PNAE)**[25], which focuses on the collection of participating entities. Pane provides serves analog to the system calls to networking systems, which allow the participating entities to provides hints to network systems and make changes to benefit their needs within allowed privileges. PANE' structure nicely forms the base of our allocation mechanism. None of the systems aforementioned utilize a central view of the network, thus we argue that our approach comes with naturally more potential. In fact, only our system supports seamless switch between strategies and this feature comes naturally with the sets up.

4 CONTRIBUTION

Our contribution are mainly building a marketing prototype over SDN and reaching some sort of marketing

fairness by our allocating algorithm . Also, our system could reach the performance requirement of real-time allocation, and supports further extension.

4.1 Floodlight and SDN based

Our work is mounted on Floodlight as its basement. With relatively mutual developed system, Floodlight and can be utilized in real world with few set ups thus make deployment of our project easier.

4.2 Marketing Fairness

Bob Briscoe in his paper[26] has bluntly pointed out that many fairness based on flows are pointless. He argues that flows are simply not the right entities to provide fairness to. We believe in his argument, but it is also arguable that sometimes fairness over flows can have its usages. A great benefits that we obtain from utilizing auction based Network resource allocation is that the fairness issue is no longer so relevant. It is up to the user to reserve and use the network resources. The Auctioneer treats everyone equally as they all have to bid.

4.3 Solve Fast

Different from allocation for computing resources, the allocation for Network resources requires a much high decision speed. In [19], Sai proposed the pattern for agents to keep bidding and potentially face the failure. This works in computation and storage allocation but is probably not the best strategy in the design of our system. In our work, other than the adjustable time needed to gather the requests for a bidding round, our controller can choose to utilize algorithms and strategies that compute fast and obtain a solution that may not be optimal but good enough. The Income-Utilization Strategy we implemented is $o(nlogn)$ in time complexity and linear in space thus can produce allocation really fast. Some of the computation such as latency validation of requests can even be done in parallel while the system is gathering bids. Thus we argue that given proper strategies, our system is scalable and fast.

4.4 Extendable Strategies

Strategies are nothing special but plug-gable modules that can be switched immediately. Although few strategies have been implemented, Our system could potentially support multiple optimizing modes including : "Short job first". "Maximun profits", "best utilization", "most met deadlines" and so on. It is up to the system manager to decide what to use.

5 SYSTEM DESIGN AND ALGORITHMS

5.1 Pre-Assumption

In order to let the system work, we have some pre-assumptions so as to limit the environment. First, we don't specify the unit of currency. Although in practice the type of currency should be figured out whether real or artificial, in this system prototype we assumed the we present money as a LONG type of value. Secondly, we don't care about where and how users gain their money they use to bid. Although in some existing systems like [11], all bidders are allocated money at first and then their wealth are adjusted by the bank profit and social welfare mechanism, our scenario does not have an 'invisible hand' which proforms as a government. Thirdly, our system only works on networks with full deployment of SDN networks. Currently the sysetm does not support networks with multiple islands, some of which are Open-Flow enabled while others are not. Fourthly, we only constrain users to bid for the machine they reside. So if a user in using a machine, he could only bid for flows starting from the using machine. Although supporting users to bid for other machines are not hard to achieve, we just want to limit the rights the bidders could hold in the network. Another thing is that the bidder coud only bid for the future. If it bids for the past, the time will be invalid. The final major assumption we make is that the working controller never goes down. Since we store information in main memory, if the controller crashes, all bidding information maintained in the bidder will be lost, and without the replication mechanism it will be a huge problem in our system. For now we temporarily put aside this consideration, and we plan to add the fail-over mechanism in the future.

5.2 Design and Architecture

In this section we introduce the design and architecture of our system, we rule a protocol of exchanging bidding messages which runs in the application layer. There are two major components of the system, one

is running inside the Floodlight controller, which acts as the auctioneer and the other is running on the client machine, which performs as a bidding agent for the client. These two sides communicate using the bidding protocol.
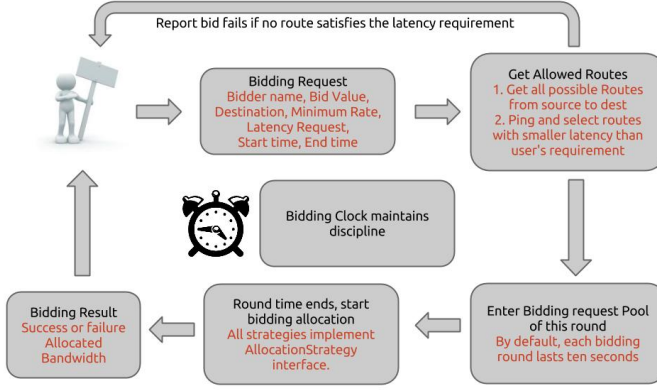

Figure 2: Bidder-Auctioneer Interaction

The interaction between bidders and auctioneers is shown in Figure 2. The bidder sends a bidding request in each round, currently, the system needs the user to fill in required fields, including his name, his bidding value, the destination of the flow he wants to bid, the minimum transform rate he accepts (in the unit of MB/s), the latency he accepts for a packet to travel from the source to the destination and the start and end time of the transmission. If some fields are what the bidder does not care about, for example the latency, he could just set a special value such as Float.MAX_VALUE to that field. But the start and end time must be set because the computation requires these two fields to be definitely set. Then the bidding request is set to the auctioneer (aka. the bidding manager). Whenever the bidding manager gets a bidding request, it immediately verifies whether the latency could be satisfied. Because the bidding manager resides in the Floodlight controller with the global topology view of the network, so it could get all possible routes using the source and destination node. Then for all possible routes, it sends probe packets and count for the time the packets use to pass all routes. The administrator could set how many times probe packets are sent for each verification in order to get an average real-time latency for the route, or the minimum if we are being rigid. If the users' required latency is smaller than the actual probed latency, we just send an bidding feedback to the client saying the bid could not be satisfied because of the latency issue. Then for each bidding request, if one of its possible path could satisfy the latency requirement, these verified routes, along with the original bidding information, are packed together to get into the bidding request pool, waiting for the bidding round to come. Then for a certain period of time, the auction time arrives, all bidding requests in the pool are collected and calculated for their results. Then the results are sent to the bidders, which marks the end of the round.

A special feature of our system is that we never over saturate any links. This means we simply do not allocate more than the network can accept, at the granularity of each switch port. The bidders can have to send the packets through a rate limiting box either in their bidding system or implemented as rules in the first switch that connected to that host. Given this design, it is nature that the latency will be overall consistent regardless of the actual traffic load. While this design also requests the strategies to be smarter to allocate more while avoid saturate the bottle neck.
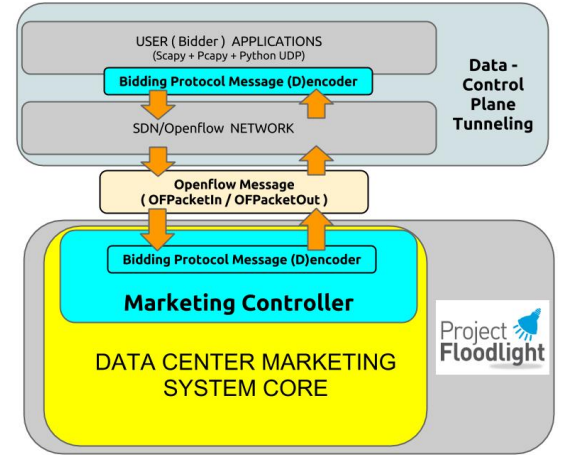

Figure 3: System Architecture

The system architecture is shown in Figure 3. The user side sends and receives bidding messages formatting in the bidding protocol. These messages are transmitted via the SDN network. As there is no control plane links from hosts to the controller. The packets are sent in the data plane to the switches, destined to the controller which resides in the control plane. We utilize the feature of OpenFlow that if the switch does not know the destination IP of a packet or it has not a rule matching the head of the packet, the swtich will transfer the packet to the controller

as a PACKET_IN message, hence letting the packet go into the control plane. So the marketing controller is listening the packets as a module of the floodlight framework. If it gets a packet from a special IP we make as a mark for the bidding message, the Marketing Controller will decode the message and packet it into a bidding request object and pass it to the Marketing Manager which lays in to logic layer, then the Manager could do the thing described in the last paragraph.

5.2.1 The Bidding Protocol

There are generally four types of packets we use to pass mesages between bidders and auctioneer. In a typical data-centers, we assume the IP addresses are highly clustered and irregular IPs does not appear. So first we assigns speicial IPs for the bidding protocol packets so that either side could recogize the special packets in the data plane. Then we serialize the packet informaiton in JSON. The reason we use JSON is that we want the protocol to be acoomodated to various machine and data presentation types.

The bidding request packet is used for the bidders to send their requests to the auctioneer. The bidder name is a String type in Java, and other values are all Long (32 bits). Since we only allows users to bid for a future period of time, the Start and End time fields in the packet must be positive values and they are presented as the relative offset of time comparing with the current time (presented as a long value in the central controller and we assume the user could know the current time). The speicial source IP for the bidding request packet is 1.2.3.4. The bidding result packet is sent by the auctioneer to the bidders. The speicial source IP is 1.2.3.4 and the bidding result is a boolean type and the allowed rate is a allocated rate which must be no smaller than the user's requested rate. The latency probe packet is used for the controller to probe the real-time latency between two nodes. The source and destination IP are set by the controller as the real IP of the two nodes (before sending the latency probe packets, the controller should establish rules on all switches on the path, using the two IP as matching fields, the out port for the last swtich is 6632 which directs the packet back to the controller). The Send Reminder Packet is used for reminding bidders to send their flows when their winning bidding time comes. The speicial source IP is 1.2.3.5. The packet tells the controller where is the destination of the flow and the duration of this transmission. The reminding mecha-

nism is facilitated by the timers so that the bidders don't need to record what they have won. The Marketing Manager records these information for the bidder and when the start time of their winning bids arrives, the reminder packet is sent. When the bidder receives the reminder packet, it will start sending UDP/TCP packets. The bidder does not need to set the route and rate limiters because the controller has already set the rules in switches on the route and set a rate limiter on one queue of the first swtich of the flow's path.

Bidding Request Packet

Special IP Header	JSON Header	Bidder Name	Bidding Value
Source Host ID	Destination Host ID	Min Rate Requirement	Data Amount
Start Time	End Time	Latency requirement	JSON Terminator

Bidding Result Packet

Special IP Header	JSON header	Bidding Result	Allowed Rate	JSON Terminator

Latency Probe Packet

Special IP Header	JSON header	Boolean Whether delete rules	Route ByteString	JSON Terminator

Send Reminder Packet

Special IP Header	JSON header	Destination	Start/End time	JSON Terminator

Figure 4: The bidding protocol
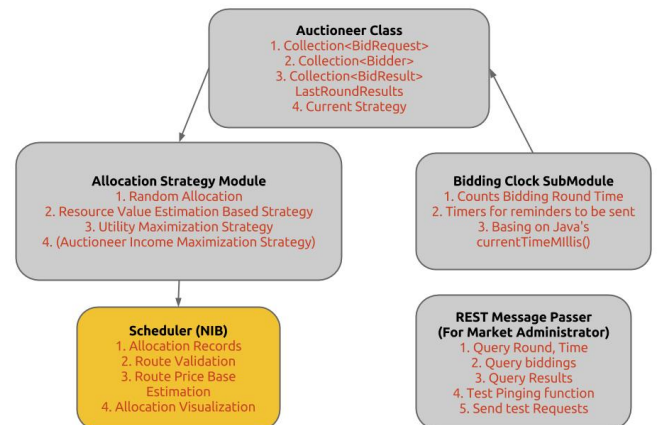
5.2.2 Marketing Manager

Figure 5: The Marketing Manager (Auctioneer) Modules

The internal modules of the Marketing Manager is shown in Figure 5. The marketing manager receives

and sends bidding related objects passed by or to the lower level bidding controller. The Auctionner object is a singleton, which maintains all bidding requests of this round (it also has a list of coming bidding requests of next round in case it is calculating the allocation for this round but some bidders have started bidding for the next round), the information of all bidders who have participated in the bidding from the start, and the results of last round. It also maintains the current allocation strategy. The bidding clock is in pace of the real world time. It sends a signal every 10 seconds (the bidding cycle time we set) to the auctioneer to perform an auction calculation. In the allocation strategy module we provide three strategies, one is first-come-first-serve random allocation, the second is the resource value estimation based strategy, which is also called Income and Utilization Index based strategy, and the third is a mutation of the previous one, which aims to maximize the utilization of bandwidth resource. A rather important part is the Scheduler module, which serves as our Network Informantion Base. The Scheduler is responsible for record all historic allocations (mainly their starting and ending time, their routes and their sending rate). The scheduler provides an interface for allocation methods to verify whether an bidding request is valid in that certain time period by doing the checking using its recorded information. Lastly, the Marketing Manager holds a seperate module for system administrators to use REST APIs to set and test some functionalities.

5.2.3 User side bidding agent

As long as the potential bidders could follow the bidding protocol, they could build their client bidding program (the bidding agent) in any language in any form. But there ar some hints for buiding the bidding agent. One is the bidders have to listen to potential incoming packets and analyze their IP address to see whether they are bidding messages. Another is that the user has to construct the bidding packet by himself, so he could manipulate all fields of the packet.

5.3 Auction Strategies

In our system, the marketing manager can utilize different strategies to selectively satisfy the bidding requests. The bidding strategies are essentially modulized parts that can be easily selected and loaded. In this section we talk about several related strategies.

5.3.1 Implemented Strategies

We have successfully implemented three strategies: first come and first serve, income and utilization index(formally called estimated index), and utilization index. Our evaluation shows that our income and utilization index can achieve significant better performance regarding the combination utilization and income comparing to the very natural first come first serve model which is sort of the standard naive philosophy to allocate network resources. Here is the algorithm we developed to provide the income-utilization index of a given route:

Data: Desired Allocation alloc which contains the bandwidth information, Route rt, which is essentially a collection of switch ports
Result: Income-utilization Index of rt
price=0;
for *each switch port p in rt* **do**
    bandwitdthPrice= alloc.bandwidth/p.leftOverBandwidth;
    queuePrice= 1/p.numOfLeftOverQueues;
    if *bandwidthPrice>queuePrice* **then**
       | price+=bandwidthPrice;
    else
       | price+=queuePrice;
    end
end
index=alloc.moneyWillingToPay/price;
return index;

Algorithm 1: Income-Utilization Index Calculation

Each routes in collected in a bidding round thus is indexed and sorted, we then go over the list from highest index to lowest index, trying to satisfy as many routes as we can in that order. Although no formal mathematical proof is provided, we take the intuition that this approach encourages users to bid for the less contingent areas of the network at a granularity of switch ports, but also allows user to pay more and make reservation on bottleneck links. Our later evaluation section will show that this algorithms works well in practice.

If we change the molecular in the equation of computing the Income-Utilization Index to 1, we have eliminated the affect of users' bidding price thus only focus on encouraging balanced utilization of the entire network. This strategy is indeed our utilization index implementation, which was also tested by us. However we do not provide more discussion around this approach other than the information that in our test,

solo utilization index not only fails obviously in cumulative income but also sometimes fails in optimizing overall utilization of the network. The exact reason remains un-investigated, but we argue that it may perform better in bigger topologies with more bottle links than our test topologies.

5.3.2 Envisioned Strategies

We have some envisioned strategies that we did not have the time to implemented, one of them is indeed seeking for the optimal income in a strict way, some packages like SuanShu and CPLEX could help we put it into code into near future, the algorithm is decribed as following :

Data: Bidding requests received for a typical auction round, presented by a vector R:

$$(R_1, R_2, R_3, ..., R_n)$$

Result: A corresponding vector with each member a binary value indicating for each bidding request whether or not to accept their requests. Presented as:

$$(x_1, x_2, x_3, ..., x_n)$$

for *each request in the request vector* **do**
    rank request by their starting time, get the modified reqeust vector R':

$$(R_{1'}, R_{2'}, R_{3'}, ..., R_{n'})$$

end

Create G as the container is used to store all requests with overlapped request time period
for *each request r in R'* **do**
    for *each request r' in requests other than r in R'* **do**
        if *r'.startTime < r.endTime* **then**
            group r' into r's time intersection group g
        end
    end
    add g into G
end
create a container G' for grouping requests intersected with switch-port pair in their sharing time period
for *each group g in G* **do**
    for *each request r in g* **do**
        for *each switch-port pair sp in the route of r* **do**
            if *sp's group is in G'* **then**
                add r into the sp's group
            else
                create sp's group in G' for this time period, add r
            end
        end
    end
end
Initialize the result vector RV,
Set up the optimization goal, which is to maximize:

$$V((x_1, x_2, x_3, ..., x_n) = \sum_{i=1}^{n} x_i * R_i.value$$

Set up the constrans :
for *each switch-port pair sp in each g' in G'* **do**

$$\sum_{i=1}^{n} x_{i'} * R_{i'}.min_b andwidth < sp.remainRate$$

end
Algorithm 2: Imcome Optimization Algorithm

6 IMPLEMENTATION DETAILS

Our prototype system is consisted of two programs, one resides in the FloodLight and another is the bidding agent we build in Python. There are roughly 6000 lines of Java code and 500 lines of Python code.

The low level marketing controller is a Floodlight listener listening for any PACKET_IN messages. In order to let our special bidding packet with special IP to be transferred to the listener, we modified the floodlight controller (implementation of IFloodlight-Provider) to let it not filter out the bidding mesasge packets. The route is presented by the Floodlight as a list of switch-port pairs, and Floodlight provides an interface to list all possible routes between any two switches, so an additional hashmap should be used to recordthe host machiens and their attachment points. In the client side, we use scapy to construct packets and forcely push them into the ethernet, so as to ensure the packet could be pushed to the attached switch. Then we use pcapy to sniff all packets coming and select packets with IP address 1.2.3.4 or 1.2.3.5. There are mainly two threads, one is for sending bidding requests and the other is listening to the coming packets in UDP.

7 EVALUATION

We have tested the performance with the aid of Mininet[5]. We have primarily tested First Come First Serve, Income and Utilization Index over three different topologies for more than 20 bidding rounds each. As an overview, our Income-Utilization index helps to achieve similar or a little better overall utilization, while achieves significant better income/utilization rate. Figures 6-14 show the details.


Figure 6: Line Topology
"Mr. LAN"

8 DISCUSSION AND FUTURE WORK

Our system is a good starting prototype, the novel idea was there but lots of potential components are remaining to be further developed and researched. Some particular interesting potential future work including :


Figure 7: Circle Topology
"The Ring"

Figure 8: Complex Topology
We call it complex because it contains a circle and a line.

8.1 Stability Issues

A significant unsatisfactory we have so far is the stability issue when we test on our Mininet testbed. The stability issues happens in multiple forms. One clear understood common reason the system fails is when the user started to bid before the internal controller configure everything it needs. The coming in bidding packet could result in exceptions in various place of the system, depending the stage the system is currently within intiliazation phase. Besides the obvious easy to fix issues, we sometimes experience random time out of virtual switches simulated by mininet. A very common issue is that, floodlight controller would suddenly timeout the netty channel

Figure 9: The utilization comparison in circle topology


Figure 10: The utilization comparison in line topology

Figure 11: The utilization comparison in complex topology
The aggressiveness of FCFS is very obvious, and FCFS appreantly does not do well regarding the long run.

through which a switch is connected, due to read time out. Almost instantly after the switch was timed out


Figure 12: The profit (income) comparison in line topology
The lines crawling underneath are for income per round.

Figure 13: The profit (income) comparison in circle topology
The lines crawling underneath are for income per round.

and disconnected, the controller will detect a new switch in mininet, with different real machine port number. This issue happens occasionally, its reason is unknown yet and could resides in at least one of floodlight, netty framwork, mininet or the reference swicth implementation.

8.2 More Tests and immigrate to real world test bed

Because of the limit of time and the delay caused by bugs hard to trace(ie, the switch loss bug mentioned aforehead), we did not have time to test our system with real switches. It is definitely an item into our todo list and needs to be done in the future.

11

Figure 14: The profit (income) comparison in complex topology
The lines crawling underneath are for income per round.

8.3 Support of Partial Bidding Request

In our current settings, users can make request without any requirements on Latency. The user could simply indicate that the latency is not important and the bidding agent could translate as and Long.MAX_VALUE to put into the request message. However, we have not provided a standard when the client only have request for latency but not bandwidth. As a bandwidth of zero make no sense regardless of its associated latency, this problem may further redirected to be to help user determine their needs for bandwidth.

8.4 Request Translation

Another very realistic request from users may be the following. Most of the users are not able to provide an information detailed enough to assemble a bid request for our system. For example, user Bob may only want to download this one Gigabyte movie tonight so he could watch it tomorrow, thus the latency is not so relevant and the shape of his allocation is elastic and dividable. Or Alice may want to schedule a phone call after five minutes and does not really want to specify an end time for it, as she does not know in advance. Her request might be composed of an minimun latency, an minimum bandwidth in an extendable session. All those real world requests need to be translated by our bidding agent to fits in our bidding system, and our bidding system could extend its rules to fit more complex and realistic scenarios.

8.5 Request Predication

Yet the biggest gap between our system to widely pratical usage is that most of the regular users may not want to make a bid, either not able to or no bother to. Thus a great break through that will benefit our system greatly would be enable our bidding agent to bid for the users by predicting the potential requests. Prediction of user requests is solo a much harder problem that people are trying to tackle(example ref paper), and its advance could benefit not only our datacenters marketing system but also many other existing systems. But as aforementioned, it is generally very hard to achieve in current state of art.

9 RELATED WORK AND COMPARISON

Data-centers have rapidly evolved to become the dominating server style in almost any modern Internet based organizations. As the amount of computation and associated network traffic keeps increasing, the needs for efficient and fair allocation in Network resources has also becomes a important topic. There exist enormous number of devices with all kinds of different hardware and software among all the datacenters. Different processes and users could have different requests and requirements. A Hadoop cluster running Map-reduce computations may requires great bandwidth. A web search engine clusters such as Google may want to response to a request within a relatively rapid time. All of the needs are to be solved with appropriate network resources allocation.

Allocating **computing and storage** resources such as Cpus and Memory in clusters and date-centers has been a popular topic widely discussed. Researchers have proposed and implemented various algorithms to achieve fairness and efficiency among different entities. Rajkumar Buyya el at' [16] proposed architecture for market-oriented allocation of resources within Clouds that encompass both customer-driven service management and computational risk management to sustain SLA-oriented resource allocation.

Artur Czumaj el at'[17] proposed the first thorough theoretical study of the price of selfish routing in server farms for general cost function, giving the hypothesis that distributed entities in data-centers are selfishly motivated. Ali Ghodsi el at[18] presented fair allocation regarding dominant resources of multiple types in a data center clusters. Sai Rahul Reddy

12

P[19] in his Master's thesis proposed combined time and budget optimized auction algorithms in grid computing.

However, there exists very few research focusing on the allocation of **network** resources. We believe the allocation on **network** resources such as guaranteed bandwidth and latency are as important as the allocation of computing and memory resources. Computation and storage are not free thus are entitled to allocation based on market mechanism, so is network resources. This is actually intuitive and commonly accepted patterns in real life. Phone carries would sell you a phone for very cheap price and make benefits from cellular services, this is a typical example in which the importance of allocation of network resource exceed allocation of computation and storage. In data centers, inappropriate abusing of network resources can results in the failure of the whole system. It has noticed by network administrators that naively running TCP between hosts can leads to huge overhead on lost packets. Domain specific algorithms like DCTCP[20] are proposed and applied in data-centers to effectively avoid collapse and keep traffic moving. But a more agile control was still very hard to reach due to the distributed nature of inter connected systems, which lacks of a central controller with a overview for the entire network and the needs of each entities are not really be collected and analyzed to form a strategy. The entire network also do not seek to optimize anything in general. It would be great if we can have a central view of the network which contact with each participating hosts in the network and allocate resources to them efficiently. Note that allocation of network resources often requires orders of magnitude faster than allocation resources in Mapreduce style systems due to the fast nature of network transmission. Thus regarding the allocation resource type, our system is very different but necessary.

10 CONCLUSION

This project is so awesome, shu and ray are so great, please give us A. thanks. Just kidding.

11 ACKNOWLEDGMENT

References

[1] http://aws.amazon.com.

[2] http://www.projectfloodlight.org/floodlight.

[3] http://aws.amazon.com/ec2/pricing.

[4] http://aws.amazon.com/ec2/spot-instances.

[5] http://mininet.org.

[6] Albert Greenberg and Gisli Hjalmtysson and David A. Maltz and Andy Myers and Jennifer Rexford and Geoffrey Xie and Hong Yan and Jibin Zhan and Hui Zhang, "ABSTRACT A Clean Slate 4D Approach to Network Control and Management"

[7] Koponen, Teemu, et al. "Onix: a distributed control platform for large-scale production networks." Proceedings of the 9th USENIX conference on Operating systems design and implementation. USENIX Association, 2010.

[8] McKeown, Nick, et al. "OpenFlow: enabling innovation in campus networks." ACM SIGCOMM Computer Communication Review 38.2 (2008): 69-74.

[9] Stokely, Murray, et al. "Using a market economy to provision compute resources across planet-wide clusters." Parallel & Distributed Processing, 2009. IPDPS 2009. IEEE International Symposium on. IEEE, 2009.

[10] Nisan, Noam. "Bidding and allocation in combinatorial auctions." Proceedings of the 2nd ACM conference on Electronic commerce. ACM, 2000.

[11] AuYoung, Alvin, et al. Practical market-based resource allocation. University of California, San Diego, 2010.

[12] Bredin, Jonathan, David Kotz, and Daniela Rus. "Market-based resource control for mobile agents." Proceedings of the second international conference on Autonomous agents. ACM, 1998.

[13] Vickrey, William. "Counterspeculation, auctions, and competitive sealed tenders." The Journal of finance 16.1 (1961): 8-37.

[14] Ausubel, Lawrence M., and Peter Cramton. "Auctioning many divisible goods." Journal of the European Economic Association 2.23 (2004): 480-493.

[15] Cramton, Peter, and Lawrence M. Ausubel. "The clock-proxy auction: A practical combinatorial auction design." (2006): 115-138.

[16] Buyya, Rajkumar, Chee Shin Yeo, and Srikumar Venugopal. "Market-oriented cloud computing: Vision, hype, and reality for delivering it services as computing utilities." High Performance Computing and Communications, 2008. HPCC'08. 10th IEEE International Conference on. Ieee, 2008.

[17] Czumaj, Artur, Piotr Krysta, and Berthold Vcking. "Selfish traffic allocation for server farms." Proceedings of the thiry-fourth annual ACM symposium on Theory of computing. ACM, 2002.

[18] Ghodsi, Ali, et al. "Dominant resource fairness: fair allocation of multiple resource types." USENIX NSDI. 2011.

[19] Reddy, Sai Rahul. "Market economy based resource allocation in grids." Master's thesis, Indian Institute of Technology, Kharagpur, India (2006).

[20] Alizadeh, Mohammad, et al. "Data center tcp (dctcp)." ACM SIGCOMM Computer Communication Review 40.4 (2010): 63-74.

[21] McKeown, Nick. "Software-defined networking." INFOCOM keynote talk, Apr (2009).

[22] Gude, Natasha, et al. "NOX: towards an operating system for networks." ACM SIGCOMM Computer Communication Review 38.3 (2008): 105-110.

[23] Foster, Nate, et al. "Frenetic: A network programming language." ACM SIGPLAN Notices 46.9 (2011): 279-291.

[24] Voellmy, Andreas, and Paul Hudak. "Nettle: Taking the sting out of programming network routers." Practical Aspects of Declarative Languages. Springer Berlin Heidelberg, 2011. 235-249.

[25] Ferguson, Andrew D., et al. "Participatory networking." Proc. Hot-ICE 12 (2012).

[26] fairness paper